



An Easy and Flexible Deep Learning Framework for PyTorch

Frédéric Paradis^{*}, David Beauchemin, Mathieu Godbout, François Lavolette

Department of Computer Science and Software Engineering, Université Laval

^{*}frederik.paradis.1@ulaval.ca

June 9, 2020



Group for
Research in
Artificial
Intelligence of
Laval University



UNIVERSITÉ
LAVAL

What Is Poutyne?

- Framework for training neural networks with PyTorch
- Includes checkpointing and logging mechanisms
- Allows to setup experiments quickly

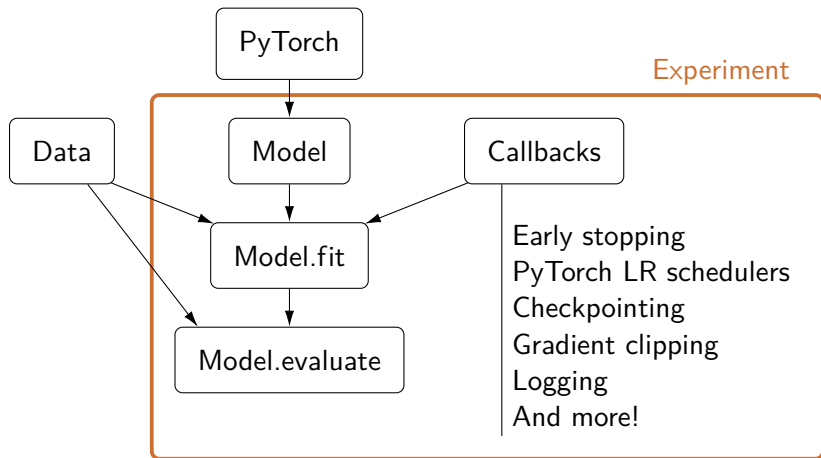


Core Principles

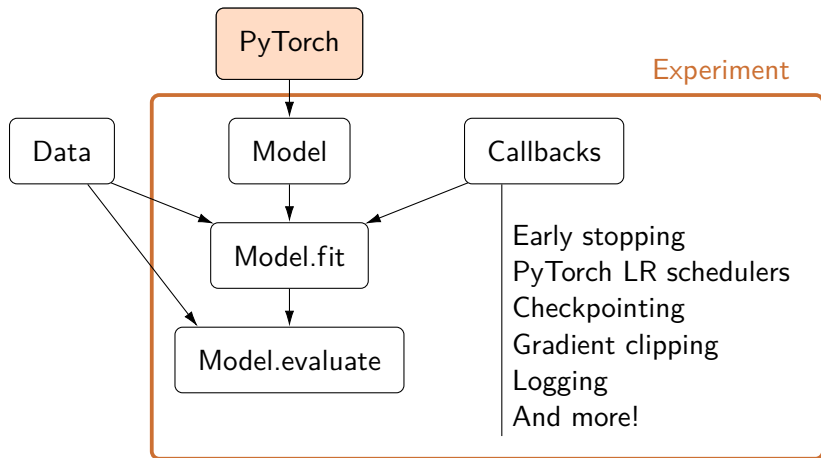
- Easy to use for simple use cases
- Flexible enough for more complex use cases
- Callbacks are your friends



Poutyne Flow



Poutyne Flow



- Automatic differentiation Python library
- For every differentiable operation done in the “forward” pass, backpropagation is done in the “backward” pass.

```
from torch import nn

class MnistLogistic(nn.Module):
    def __init__(self):
        super().__init__()
        self.weights = nn.Parameter(torch.randn(784, 10) /
                                         math.sqrt(784)) # W
        self.bias = nn.Parameter(torch.zeros(10)) # b

    def forward(self, xb):
        return xb.matmul(self.weights) + self.bias #  $xW + b$ 
```



Usual Code for Training With PyTorch

```
import torch
import torch.nn as nn
import torch.optim as optim
from torch.utils.data import TensorDataset, DataLoader

net = nn.Sequential(
    nn.Linear(100, 64),
    nn.ReLU(),
    nn.Linear(64, 10)
)

num_features = 100
num_classes = 10

train_dataset = TensorDataset(x_train, y_train)
train_loader = DataLoader(train_dataset, batch_size=32)

test_dataset = TensorDataset(x_test, y_test)
test_loader = DataLoader(test_dataset, batch_size=32)

criterion = nn.CrossEntropyLoss()
optimizer = optim.SGD(net.parameters(), lr=0.001, momentum=0.9)
```

```
for epoch in range(5): # loop over the dataset multiple times
    running_loss = 0.0
    for i, data in enumerate(train_loader, 0):
        # get the inputs; data is a list of [inputs, labels]
        inputs, labels = data

        # zero the parameter gradients
        optimizer.zero_grad()

        # forward + backward + optimize
        outputs = net(inputs)
        loss = criterion(outputs, labels)
        loss.backward()
        optimizer.step()

        # print statistics
        running_loss += loss.item()
    if i % 2000 == 1999: # print every 2000 mini-batches
        print('%d, %5d] loss: %.3f' % (epoch + 1, i + 1, running_loss / 2000))
        running_loss = 0.0

correct = 0
total = 0
with torch.no_grad():
    for data in test_loader:
        images, labels = data
        outputs = net(images)
        _, predicted = torch.max(outputs.data, 1)
        total += labels.size(0)
        correct += (predicted == labels).sum().item()

print('Accuracy of the network on the 10000 test images: %d %%' % (100 * correct / total))
```



Equivalent Code for Training With Poutyne

```
import torch.nn as nn
import torch.optim as optim
from poutyne.framework import Model

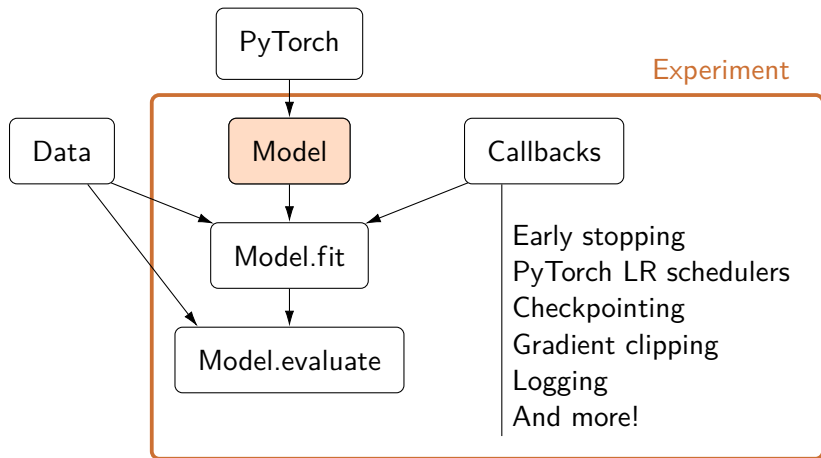
net = nn.Sequential(
    nn.Linear(100, 64),
    nn.ReLU(),
    nn.Linear(64, 10)
)

criterion = nn.CrossEntropyLoss()
optimizer = optim.SGD(net.parameters(), lr=0.001, momentum=0.9)

model = Model(net, optimizer, criterion, batch_metrics=['accuracy'])
model.fit(x_train, y_train, epochs=5, batch_size=32)
loss, accuracy = model.evaluate(x_test, y_test, batch_size=32)
```



Poutyne Flow



Model Class

- Main class of the framework
- Plays well with Numpy
- No restrictions on the input or the output format of the network
- Manages devices (GPUs)

```
from poutyne.framework import Model

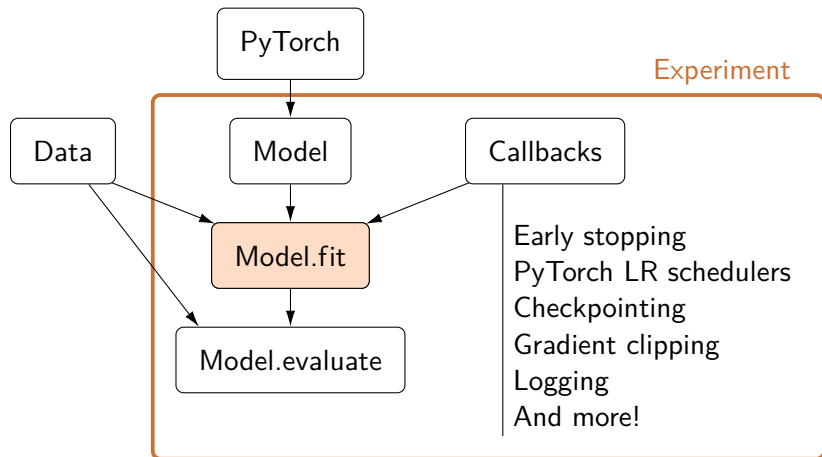
model = Model(network, optimizer, loss_function)
model.to(device)

model.fit_generator(train_loader, valid_loader,
                   epochs=num_epochs, callbacks=callbacks)


test_loss = model.evaluate_generator(test_loader)
```



Poutyne Flow



Poutyne Training Flow



Poutyne Training Flow

for n epochs **do**

for each drawn batch (x, y) in training dataset **do**

$$\hat{y} = f(x; \theta)$$

$$\ell = \mathcal{L}(\hat{y}, y)$$

$$g = \nabla_{\theta} \ell$$

Update θ with g using chosen optimizer.

Compute and accumulate metrics with \hat{y} and y .

end for

Compute loss and metrics on validation dataset.

end for



Poutyne Training Flow

for n epochs **do**

Callback on epoch begin

for each drawn batch (x, y) in training dataset **do**

Callback on batch begin

$$\hat{y} = f(x; \theta)$$

$$\ell = \mathcal{L}(\hat{y}, y)$$

$$g = \nabla_{\theta} \ell$$

Callback on backward end

 Update θ with g using chosen optimizer.

 Compute and accumulate metrics with \hat{y} and y .

Callback on batch end

end for


 Compute loss and metrics on validation dataset.

Callback on epoch end

end for



Poutyne Flow



Evaluation Metrics

Using metrics with Poutyne allows easy early stopping, checkpointing and logging.



Evaluation Metrics

Using metrics with Poutyne allows easy early stopping, checkpointing and logging.

- Batch metrics
 - Decomposable metrics (e.g. accuracy, mse, etc.)
 - Any PyTorch loss function



Evaluation Metrics

Using metrics with Poutyne allows easy early stopping, checkpointing and logging.

- Batch metrics
 - Decomposable metrics (e.g. accuracy, mse, etc.)
 - Any PyTorch loss function
- Epoch metrics
 - Non-decomposable metrics (e.g. F1, AUC ROC, etc.)
 - Provide a scikit-learn wrapper



Evaluation Metrics

Using metrics with Poutyne allows easy early stopping, checkpointing and logging.

- Batch metrics
 - Decomposable metrics (e.g. accuracy, mse, etc.)
 - Any PyTorch loss function
- Epoch metrics
 - Non-decomposable metrics (e.g. F1, AUC ROC, etc.)
 - Provide a scikit-learn wrapper

```
from poutyne.framework import Model

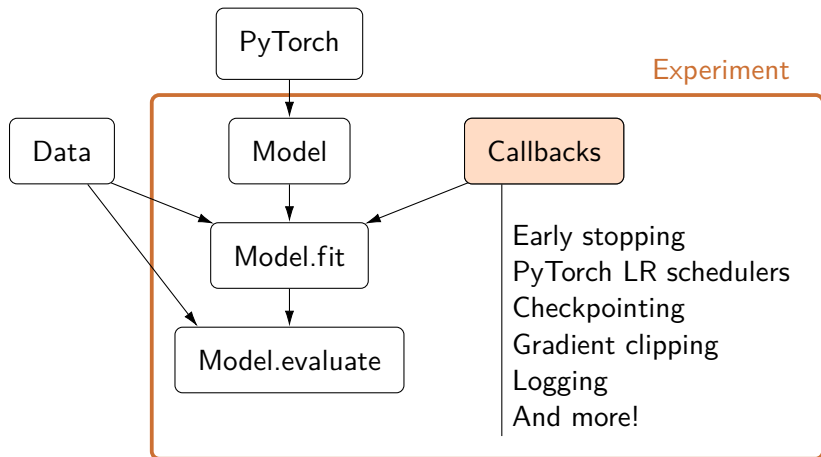
model = Model(network, 'sgd', 'cross_entropy',
               batch_metrics=['accuracy'], epoch_metrics=['f1'])
model.to(device)

model.fit_generator(train_loader, valid_loader,
                   epochs=num_epochs, callbacks=callbacks)

test_loss, (test_acc, test_f1) = model.evaluate_generator(test_loader)
print(f'Test: Loss: {test_loss}, Accuracy: {test_acc}, F1: {test_f1}')
```



Poutyne Flow



Callbacks

```
class Callback:
    def on_train_begin(self, logs: dict): ...
    def on_train_end(self, logs: dict): ...

    def on_epoch_begin(self, epoch_number: int, logs: dict): ...
    def on_epoch_end(self, epoch_number: int, logs: dict): ...

    def on_train_batch_begin(self, batch_number: int, logs: dict): ...
    def on_train_batch_end(self, batch_number: int, logs: dict): ...

    def on_backward_end(self, batch_number: int): ...

    def on_test_batch_begin(self, batch_number: int, logs: dict): ...
    def on_test_batch_end(self, batch_number: int, logs: dict): ...

    def on_test_begin(self, logs: dict): ...
    def on_test_end(self, logs: dict): ...

    self.params = {...} # Contains 'epochs' and 'steps_per_epoch'
    self.model = ... # Poutyne Model
```



Callbacks

```
from poutyne.framework import Model, ModelCheckpoint, CSVLogger

callbacks = [
    ModelCheckpoint('last_epoch.ckpt'),
    ModelCheckpoint('best_epoch.ckpt', save_best_only=True,
                    monitor='val_acc', mode='max'),
    CSVLogger('log.csv'),
]

model = Model(network, 'sgd', 'cross_entropy',
               batch_metrics=['accuracy'], epoch_metrics=['f1'])
model.to(device)

model.fit_generator(train_loader, valid_loader,
                   epochs=num_epochs, callbacks=callbacks)

test_loss, (test_acc, test_f1) = model.evaluate_generator(test_loader)
print(f'Test: Loss: {test_loss}, Accuracy: {test_acc}, F1: {test_f1}')
```



Checkpointing

- ModelCheckpoint
- OptimizerCheckpoint
- LRSchedulerCheckpoint



Early Stopping and LR Scheduling

- EarlyStopping
- Any PyTorch LR scheduler
- FastAI-like learning rate policies




Playing With Gradient

- Use PyTorch's respective functions
 - ClipNorm
 - ClipGrad
- TensorBoardGradientTracker



Poutyne Flow



Experiment Class

- Allows stopping and resuming optimization any time
- Saves and logs everything in a single directory
- Integration with Tensorboard

```
from poutyne.framework import Experiment

# Instead of `task`, you can provide your own loss function and metrics.
expt = Experiment('my_directory', network,
                  task='classifier', optimizer='sgd')
expt.train(train_loader, valid_loader,
           epochs=epochs,
           callbacks=callbacks,
           seed=42)
expt.test(test_loader)
```



Experiment Class

- Saves the last checkpoint and every “best” checkpoint (ModelCheckpoint).
- Saves the last states of the optimizer and LR schedulers (OptimizerCheckpoint, LRSchedulerCheckpoint).
- Logs training and validation loss and metrics into CSV and Tensorboard (CSVLogger, TensorBoardLogger).



Related Works

PyTorch Lightning:

- Flexible
- Couples network with training
 - Requires inheriting from a special class
 - Everything training related is inside the network
- Add boilerplate where it should not (e.g. adding LR schedulers seems awkward)



Related Works

PyTorch Lightning:

- Flexible
- Couples network with training
 - Requires inheriting from a special class
 - Everything training related is inside the network
- Add boilerplate where it should not (e.g. adding LR schedulers seems awkward)

FastAI:

- Lots of features
- API not as intuitive



Related Works

PyTorch Lightning:

- Flexible
- Couples network with training
 - Requires inheriting from a special class
 - Everything training related is inside the network
- Add boilerplate where it should not (e.g. adding LR schedulers seems awkward)

FastAI:

- Lots of features
- API not as intuitive

AllenNLP:

- Specialized for NLP
- Experiment framework
- Trainer not flexible enough (everything in `__init__`)



Demo Time

Fine-tuning with dataset:

<http://www.vision.caltech.edu/visipedia/CUB-200.html>

Future Works


- Add tqdm and colors for progression.
- Add tutorial pages to website.
- Integrate multi-GPU.
- Add a simpler way to add regularizer to the loss function.
- Add utilities to simplify parameters initialization.
- Integrate an experiment library such as MLFlow.




Obtain Poutyne




Install via pip

 `pip install poutyne`

Documentation and examples available!

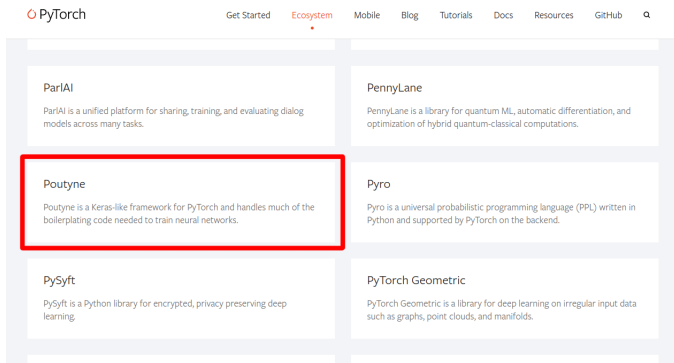
 <https://poutyne.org>

 [GRAAL-Research/poutyne](https://github.com/GRAAL-Research/poutyne)



Poutyne in the PyTorch Ecosystem

<https://pytorch.org/ecosystem/>



The screenshot shows the PyTorch Ecosystem page. The header includes the PyTorch logo and navigation links: Get Started, Ecosystem (highlighted with a red dot), Mobile, Blog, Tutorials, Docs, Resources, GitHub, and a search icon. The main content area is a grid of project cards. The 'Poutyne' card is highlighted with a red rectangular border. The card contains the title 'Poutyne' and a description: 'Poutyne is a Keras-like framework for PyTorch and handles much of the boilerplating code needed to train neural networks.'

PyTorch

Get Started **Ecosystem** Mobile Blog Tutorials Docs Resources GitHub

ParlAI
ParlAI is a unified platform for sharing, training, and evaluating dialog models across many tasks.

PennyLane
PennyLane is a library for quantum ML, automatic differentiation, and optimization of hybrid quantum-classical computations.

Poutyne
Poutyne is a Keras-like framework for PyTorch and handles much of the boilerplating code needed to train neural networks.

Pyro
Pyro is a universal probabilistic programming language (PPL) written in Python and supported by PyTorch on the backend.

PySyft
PySyft is a Python library for encrypted, privacy preserving deep learning.

PyTorch Geometric
PyTorch Geometric is a library for deep learning on irregular input data such as graphs, point clouds, and manifolds.



The end.

Questions?

